How to write a rebuttal

ANDREW FITZGIBBON

My CV

About

Bio

Andrew Fitzgibbon leads the "All Data AI" (ADA) research group at Microsoft in Cambridge, UK.

He is best known for his work on 3D vision, having been a core contributor to the Emmy-award-winning 3D camera tracker "boujou", to body tracking for Kinect for Xbox 360, and for the articulated hand-tracking interface to Microsoft's HoloLens.

His research interests are broad, spanning computer vision, machine learning, programming languages, computer graphics and occasionally a little neuroscience.

He has published numerous highly-cited papers, and received many awards for his work, including ten "best paper" prizes at various venues, the Silver medal of the Royal Academy of Engineering, and the BCS Roger Needham award. He is a fellow of the Royal Academy of Engineering, the British Computer Society, and the International Association for Pattern Recognition, and is a Distinguished Fellow of the British Machine Vision Association.

Before joining Microsoft in 2005,...

Read more﹀

My CV – relevant to this workshop…

About

Bio

Andrew Fitzgibbon leads the "All Data

He is best known for his work on 3D vision, having been a core contributor to the Emmy-award-winning 3D camera tracker "boujou", to body tracking for Kinect for Xbox 360, and for the articulated hand-tracking interface to Microsoft's HoloL

His research interests are broad, spann computer graphics and occasionally a

He has published numerous highly-cite "best paper" prizes at various venues, t Roger Needham award. He is a fellow and the International Association for P Machine Vision Association.

Before joining Microsoft in 2005,...

Read more⌄

```
%
% $Id: cvpr.sty,v 1.3 2005/10/24 19:56:15 awf Exp $
%
% by Paolo.Ienne@di.epfl.ch
% some mods by awf@acm.org
%
```

are known and available.

Blind review means that you do not use the words "my" or "our" when citing previous work. That is all. (But see below for techreports.)

Saying "this builds on the work of Lucy Smith [1]" does not say that you are Lucy Smith; it says that you are building on her work. If you are Smith and Jones, do not say "as we show in [7]", say "as Smith and Jones show in [7]" and at the end of the paper, include reference 7 as you would any other cited work.

An example of a bad paper just asking to be rejected:

My experience as an author

- Most reviews have got the idea of the paper. Very few reviewers "haven't read the paper".
- Most reviews contain some misunderstanding of some part of the paper. And **this misunderstanding is mostly my fault**. I was ambiguous, I rambled, I had a mistake in notation, I didn't follow Bill's rules. The reviews can help me fix that.
- Sometimes they tell me things I really, deep down, knew. I knew we needed that baseline, I knew we really needed to do the C++ implementation to prove our claim it was fast.

My experience as an area chair

- Most reviewers have got the idea of the paper. Very few reviewers "haven't read the paper". Those who haven't are very easy to ignore.
- Some reviewers have a "bee in their bonnet" and focus on something unimportant. As an AC I know how to balance it.
- Some reviewers point out possible prior art [23]. As an AC, I can carefully read [23], but if the reviewer pointed out the precise equation in their prior art, that helps me. If they did not, **the rebuttal** can help me.
- If the reviewer's prior art was somehow wrong ([23] uses 3D data, this paper uses 2D). I *should* spot that, but if I do not, **the rebuttal** can help me.

So…

The reviews are in!

The reviews have arrived!!! I hope they're all "Strong accept"!!

They must be, right?

- We beat X on [benchmark]!
- We combined A&B to do C, and no one thought of that before.

Oh, they're not…. What now?

IDIOTS!

NINCOMPOOPS!!

BIRDBRAINS!!!!

KEEP
CALM
AND
BREATHE

Rebuttal purpose and audience

PURPOSE

To get my paper accepted/oral/prize.

Possibly: to help me understand what I can improve.

Content: Refutation of errors of misunderstanding.

AUDIENCE

Reviewers

- for discussion phase.
- you're reviewing their review

Area chair

- to seed discussions,
- and then to make a decision.

Rebuttal meta-notes:

Different venues have different rules:

- Extra experiments not allowed, allowed only if requested, encouraged, …
- N characters/words/spaces per rebuttal/reviewer/…

These notes to be taken *mutatis mutandis*

Rebuttal steps:

Gather reviews into one shared doc.

Add newlines after every comment, positive or negative.

Get calm. [Can be done in parallel with above steps]

Read the comments. Carefully. And be generous: did the reviewer mistype something?

Ask: are they right? If so, can you fix it? If not, consider withdrawing [separate slide].

Add notes – what is the essence of their complaint? What would we need to do to resolve it, given infinite time?

Sort by importance, and answer precisely, concisely, unambiguously.

Rejection

Not the worst thing in the world. Easy to say if all authors have tenure/industry jobs, for someone approaching end of PhD, it can have real and immediate financial impact.

But *there is an upside* to rejection:

- If the rebuttal had worked, it would have probably at best got you a poster.
- A resubmission which learns from the reviews might well do much better.
  (I have seen reject → Prize)
- But as Bill says, even then, a brilliant paper serves you better than an OK one. And remember, it's the *paper* that's only OK, not the idea.

Rebuttal nitty gritty:

I think it's worth a "General" section at the top, addressing any issues that apply to multiple reviewers.

"We thank all reviewers for their comments. R2 and R3 note that we did not compare to [A].
We were not aware of this work at the time of submission.
They propose a new loss, which would replace our term L_foo with EMD.
As our primary contribution is in the optimization of the loss (line 123),
we [predict that / have expts to show that] our contribution applies to [A] too.

[A] Emeagwali, Darden, CVPR17

And be precise and accurate (and of course honest).

Rebuttal nitty gritty:

And then per-reviewer sections.

For each reviewer, ask “why didn’t they love it?”

- It *might* be that they were annoyed you beat them.
- Or they’re suspicious: they don’t see why this new complex thing should work.
- Or they just went “meh”.

They may have used “excuses” to reject, but try to excite them.

X killed my paper!

- I know it was X, because s/he named a bunch of X papers I didn't cite
- I know it was X, because our results are better than theirs.
- I know it was X, because s/he is well known for hating [idea Y]

- Honestly, it really probably isn't. I've had many situations where I've been PC/AC, so known the reviewer, and had people tell me they are certain it's X. There are just too many papers, too many reviewers for X to review them all.
- I know it's cool to say "Oh, I have a feud with X", and it's much better than saying "my paper had flaws, I'm working on them"

Conclusions

- Be calm
- Be polite
- Be accurate
- If you were wrong, own it
- Know your audience (AC & Reviewers)
- Use the opportunity to learn